PROOFLINE — ROUND 2 PPT REVISION INSTRUCTIONS

Team Binary Brains • Verified Slide-by-Slide Copy Changes (Slides 1 to 10)

SLIDE 1: COVER / TITLE SLIDE

• **Title:** PROOFLINE
• **Subtitle:** Evidence-Backed Brand Governance — Verify Before You Publish
• **Presented By:** Binary Brains
• **Required Addition:** Add Live Demo Link: [Insert Live URL]
• **Required Addition:** Add GitHub Repository Link: [Insert Repo URL]

SLIDE 2: PROBLEM STATEMENT

[Sub-headline] FROM: AI-generated content is growing rapidly, but businesses still struggle to ensure that every claim follows their latest approved facts and brand guidelines.

➔ **TO: Content creation scaled 10x with AI, but businesses still struggle to verify that every claim matches their latest approved facts.**

[Section Header] FROM: WHY ARE WE DOING THIS?

➔ **TO: WHERE TRADITIONAL REVIEW BREAKS**

[Bullet 1] FROM: Content Inconsistency: Multiple creators and teams may produce content that does not follow the same brand guidelines.

➔ **TO: Disconnected Guidelines: Brand rules sit in static PDFs that creators and AI tools never check while drafting.**

[Bullet 2] FROM: Incorrect Claims: Product specifications or promotional statements may conflict with approved information.

➔ **TO: Contradicted Claims: Marketing copy frequently clashes with real specs (e.g., claiming 20h battery instead of approved 14h).**

[Bullet 3] FROM: Outdated Information: Old product details and expired offers can remain in marketing content.

➔ **TO: Expired Information: Deprecated pricing and lapsed offers slip into live campaigns without version checks.**

[Bullet 4] FROM: Manual Verification: Teams must repeatedly check documents, product information, and brand guidelines before publishing.

➔ **TO: Bottlenecked Manual Review: Reviewers waste hours repeatedly cross-checking docs instead of high-value work.**

• **Bottom Callout:** The Core Gap: Content production scaled 10x with AI, but verification is still manual, unversioned, and reactive.

SLIDE 3: TARGET USERS (WHO NEEDS PROOFLINE?)

[Sub-headline 1] FROM: WHO NEEDS PROOFLINE?

➔ **TO: WHO USES THE WORKFLOW?**

[Sub-headline 2] FROM: Different teams. One shared need: trustworthy brand content.

➔ **TO: Different roles. One shared pre-publication quality gate.**

[Bullet 1 (Creators)] FROM: CONTENT CREATORS: Verify claims and brand guidelines before submitting content.

➔ **TO: CONTENT CREATORS: Catch contradicted specs and get instant, intent-preserving corrections before submission.**

[Bullet 2 (Marketing)] FROM: MARKETING TEAMS: Maintain consistent messaging across campaigns and channels.

➔ **TO: CAMPAIGN MANAGERS: Triage flagged drafts by risk score, compare revisions, and track approval status.**

[Bullet 3 (Reviewers)] FROM: BRAND REVIEWERS: Check supporting evidence before approving content.

➔ **TO: BRAND REVIEWERS: Inspect evidence-linked sources and log traceable approvals (no blind trust in AI).**

[Bullet 4 (Admins)] FROM: BRAND ADMINISTRATORS: Manage approved facts, sources, and brand rules.

➔ **TO: FACT & BRAND OWNERS: Maintain the single source of truth—versioning specs, validity dates, and brand rules.**

• **Bottom Callout:** Target Setting: High-velocity mid-market brands, agencies, and consumer tech scaling multi-channel content.

SLIDE 4: THE SOLUTION

[Sub-headline] FROM: Evidence-Backed AI for Brand Governance

➔ **TO: A Pre-Publication Quality Gate Powered by Source-Grounded AI**

[Box 1 Title] FROM: Claim–Fact Ledger

➔ **TO: Brand Truth Registry + Claim Ledger**

[Box 1 Body] FROM: Connects every detected content claim to approved brand facts, source versions, and applicable rules.

➔ **TO: Centralises versioned brand specs and links every detected claim to its exact source document and validity date.**

[Box 2 Title] FROM: Evidence-Based Verification

➔ **TO: Deterministic & Grounded AI Auditing**

[Box 2 Body] FROM: Checks product specifications, promotional claims, brand terminology, and content against approved information.

➔ **TO: Combines exact deterministic rules (numbers, pricing, banned terms) with source-grounded LLM context checks.**

[Box 3 Title] FROM: Outdated & Unsupported Claim Detection

➔ **TO: Evidence-or-Abstain Reliability**

[Box 3 Body] FROM: Flags expired information, contradictions, and claims that lack sufficient evidence instead of guessing.

➔ **TO: Detects expired promos and contradicted specs; marks missing evidence as 'Needs Human Review' instead of hallucinating.**

[Box 4 Title] FROM: Human Review & Correction

➔ **TO: Human-in-the-Loop Review & Audit Trail**

[Box 4 Body] FROM: Suggests minimal corrections, supports rechecking revised content, and records reviewer decisions for accountability.

➔ **TO: Suggests intent-preserving fixes, supports re-auditing revised drafts, and locks final sign-off to authorised humans.**

SLIDE 5: COMPLETE PRODUCT WORKFLOW

[Step 1 Title] FROM: Brand Knowledge Setup

➔ **TO: Brand Truth Setup**

[Step 1 Subtitle] FROM: Upload approved facts, rules and source documents

➔ **TO: Register versioned specs, validity dates, and brand guidelines**

[Step 2 Title] FROM: Content Input

➔ **TO: Multi-Format Intake**

[Step 2 Subtitle] FROM: Submit draft content for verification

➔ **TO: Ingest marketing copy, PDFs, DOCX, or OCR image assets**

[Step 3 Title] FROM: Claim Extraction & Analysis

➔ **TO: Claim Extraction**

[Step 3 Subtitle] FROM: Identify claims and check against brand rules

➔ **TO: Isolate factual specs, pricing numbers, and marketing assertions**

[Step 4 Title] FROM: Evidence Matching

➔ **TO: Source-Grounded Audit**

[Step 4 Subtitle] FROM: Match claims with approved sources and flag issues

➔ **TO: Classify into Supported, Contradicted, Stale, or Unverified**

[Step 5 Title] FROM: Corrections & Recheck

➔ **TO: Correction & Re-Scan**

[Step 5 Subtitle] FROM: Suggest corrections and re-verify revised content

➔ **TO: Apply intent-preserving fixes and re-audit revised drafts**

[Step 6 Title] FROM: Human Review & Approval

➔ **TO: Authorised Sign-off**

[Step 6 Subtitle] FROM: Approve or request changes with recorded decision

➔ **TO: Review evidence diffs and lock decisions in the audit log**

[Section Header] FROM: KEY WORKFLOW CAPABILITIES:

➔ **TO: CORE WORKFLOW GUARANTEES:**

[Bullet 1] FROM: Versioned Brand Knowledge: Stores approved facts, source versions, and validity dates.

➔ **TO: Versioned Truth Registry: Tracks exact source citations, effective dates, and fact expiration windows.**

[Bullet 2] FROM: Claim Verification: Combines deterministic rules with AI-assisted content analysis.

➔ **TO: Hybrid Audit Engine: Deterministic checks for numbers/dates + LLM checks for context and tone.**

[Bullet 3] FROM: Evidence-Linked Findings: Connects detected issues to approved facts and applicable rules.

➔ **TO: Evidence-or-Abstain: Every flag cites an approved record; missing evidence flags 'Needs Review' instead of guessing.**

[Bullet 4] FROM: Revision Tracking: Records content versions and reviewer decisions.

➔ **TO: Immutable Lineage: Preserves original drafts, accepted edits, and side-by-side revision comparisons.**

[Bullet 5] FROM: Human Oversight: Keeps final approval with an authorised reviewer.

➔ **TO: Zero Auto-Publishing: Content is never published automatically; final authority stays with an authorised human.**

SLIDE 6: FEATURES & TECHNOLOGY STACK

[Slide Subtitle] FROM: Planned capabilities and implementation approach

➔ **TO: Implemented system architecture and verified build stack**

[Left Header] FROM: CORE FEATURES

➔ **TO: CORE IMPLEMENTED CAPABILITIES**

[Card 01 Body] FROM: Manage approved facts, brand rules, source versions, and validity dates.

➔ **TO: Stores versioned facts, source citations, validity windows, and brand guidelines.**

[Card 02 Title] FROM: Evidence-Backed Auditing

➔ **TO: Claim Grounding & Audit Engine**

[Card 02 Body] FROM: Detect contradicted, outdated, unsupported, and ambiguous claims.

➔ **TO: Extracts claims and audits them against sources: Supported, Contradicted, or Unverified.**

[Card 03 Title] FROM: Correction & Recheck

➔ **TO: Intent-Preserving Correction & Re-audit**

[Card 03 Body] FROM: Suggest minimal corrections and verify revised content again.

➔ **TO: Generates minimal diff-based fixes and re-audits revised drafts instantly.**

[Card 04 Title] FROM: Review & Audit Trail

➔ **TO: Human-in-the-Loop Audit Trail**

[Card 04 Body] FROM: Record approval decisions, content versions, and reviewer actions.

➔ **TO: Enforces role-based sign-off and records an immutable, timestamped review history.**

[Right Header] FROM: PROPOSED TECHNOLOGY STACK

➔ **TO: VERIFIED BUILD STACK**

[Row 1 (Frontend)] FROM: React, TypeScript, Tailwind CSS

➔ **TO: Next.js (App Router), TypeScript, Tailwind CSS**

[Row 2 (Backend)] FROM: Python, FastAPI

➔ **TO: Python 3.12, FastAPI, Pydantic v2**

[Row 3 (Database & Auth)] FROM: PostgreSQL / Supabase

➔ **TO: SQLite / D1 & Postgres schemas, scoped memory vault**

[Row 4 (AI Layer)] FROM: Configurable LLM API with structured output validation

➔ **TO: Schema-enforced LLM adapter with strict JSON output validation**

[Row 5 (Verification Engine)] FROM: Python rule-based checks + evidence retrieval

➔ **TO: Deterministic exact-match rules + source lexical retrieval**

[Row 6 (Document Processing)] FROM: PDF/DOCX text extraction; bounded image OCR (planned)

➔ **TO: Local document extraction (PDF / DOCX) + bounded OCR parser**

[Bottom Callout 1] FROM: Evidence first | Every claim backed by sources

➔ **TO: Evidence-or-Abstain | Missing evidence yields 'Needs Review', not hallucinations**

[Bottom Callout 2] FROM: AI-assisted checks | Faster and more accurate verification

➔ **TO: Deterministic-First | Numbers, pricing, and dates validated by code, not LLM guesses**

[Bottom Callout 3] FROM: Human-controlled approval | Final decision always with a reviewer

➔ **TO: Human Accountability | Zero auto-publishing; strict separation of creator and approver**

SLIDE 7: RESPONSIBLE AI & DATA SECURITY

[Bottom Banner (Typo Fix)] FROM: OUR PRINCIPLE: AI assists verification. Humans remain accountable for final approval.

➔ **TO: OUR PRINCIPLE: AI assists verification. Humans remain accountable for final approval.**

[Card 01 Title] FROM: PRIVACY BY DESIGN

➔ **TO: Tenant Isolation & Data Privacy**

[Card 01 Body] FROM: Keep uploaded content private, isolate customer data, protect API keys, and define clear data-retention and deletion policies.

➔ **TO: Enforces strict organisation workspace boundaries, keeps API secrets strictly server-side, and stores zero client keys.**

[Card 02 Title] FROM: EVIDENCE-BASED AI

➔ **TO: Evidence-or-Abstain Engine**

[Card 02 Body] FROM: Link findings to approved sources and rules. When evidence is missing, conflicting, or outdated, flag the issue for review instead of guessing.

➔ **TO: Grounds every finding in a versioned source citation; flags missing or conflicting evidence as 'Needs Review' instead of hallucinating.**

[Card 03 Title] FROM: HUMAN OVERSIGHT

➔ **TO: Accountable Human Sign-off**

[Card 03 Body] FROM: Let authorised reviewers approve content, request changes, or override findings with a recorded reason.

➔ **TO: Enforces role separation (creators cannot approve own work), logs override reasons, and strictly blocks auto-publishing.**

[Card 04 Title] FROM: TRANSPARENCY & LIMITATIONS

➔ **TO: Explicit Disclosure & Fail-Soft States**

[Card 04 Body] FROM: Show the evidence behind each finding, disclose incomplete extraction or AI failures, and avoid claiming guaranteed accuracy or legal compliance.

➔ **TO: Dedicated /privacy disclosure route, transparent parsing-failure alerts, and clear disclaimers that Proofline is governance, not legal counsel.**

SLIDE 8: DEVELOPMENT PLAN & EXPECTED IMPACT

[Box 1 Subtitle] FROM: PLAN & VALIDATE

➔ **TO: PLAN & VALIDATE (COMPLETED)**

[Box 2 Subtitle] FROM: BUILD THE MVP

➔ **TO: CORE BUILD & BASELINES (DELIVERED)**

[Box 2 Bullet 1] FROM: Develop the Brand Truth Registry and content intake.

➔ **TO: Delivered: Versioned Brand Truth Registry, validity dates, and multi-format intake.**

[Box 2 Bullet 2] FROM: Implement claim extraction, rule-based checks, and evidence-linked findings.

➔ **TO: Delivered: Hybrid audit engine (deterministic spec rules + source-grounded findings).**

[Box 2 Bullet 3] FROM: Add corrections, rechecking, and human review.

➔ **TO: Delivered: Intent-preserving diff fixes, re-audit pipeline, and role-based sign-off.**

[Box 2 Bullet 4] FROM: Test and demonstrate the end-to-end workflow.

➔ **TO: Delivered Baselines: Claim–fact lineage, explicit abstain states, and review metrics.**

[Box 3 Subtitle] FROM: IMPROVE & DEMONSTRATE

➔ **TO: EXPANDED CAPABILITIES (NEXT)**

[Box 3 Bullet 1] FROM: Add the Fact-Change Impact Radar.

➔ **TO: Full Fact-Change Impact Radar: Automated re-audit queue when specs change.**

[Box 3 Bullet 2] FROM: Strengthen Evidence-or-Abstain reliability.

➔ **TO: High-severity verifier pass and edge-case benchmark test corpus.**

[Box 3 Bullet 3] FROM: Develop Governance Insights and a Guideline-Gap Detector.

➔ **TO: Guideline-Gap Detector: Proposes rule clarifications for recurring mistakes.**

[Box 3 Bullet 4] FROM: Evaluate reliability, usability, and security.

➔ **TO: Structured user testing trials and live external provider integration.**

[Impact Pill 1] FROM: Reduce time spent checking brand claims across content.

➔ **TO: Eliminates repetitive manual spec lookups, accelerating draft-to-review velocity.**

[Impact Pill 2] FROM: Help find outdated or unsupported information before publication.

➔ **TO: Catches contradictory specs, expired promotions, and missing qualifiers pre-publication.**

[Impact Pill 3] FROM: Maintain clear records of content reviews and approval decisions.

➔ **TO: Provides an immutable audit log of source versions, diffs, and reviewer sign-offs.**

[Scope Pill 1] FROM: Connect with content-management and collaboration tools.

➔ **TO: CMS connectors (WordPress, Webflow), Google Docs add-on, and Slack notifications.**

[Scope Pill 2] FROM: Add image, multilingual, and additional channel support.

➔ **TO: Multi-image visual asset matching, video transcripts, and multi-language audits.**

[Scope Pill 3] FROM: Deeper governance insights and trend analysis.

➔ **TO: Cross-team compliance heatmaps and recurring mistake telemetry.**

[Bottom Note] FROM: Impact goals will be validated through testing and user feedback; results are not yet measured.

➔ **TO: Validation Note: Impact targets will be measured in user trials; no fabricated production metrics claimed.**

SLIDE 9: TEAM & RESPONSIBILITIES

[Subtitle] FROM: Clear ownership for building PROOFLINE

➔ **TO: Verified ownership & core engineering contributions**

[Card 01 Title] FROM: Project Lead & Research

➔ **TO: [Insert Name] — Lead & Architecture**

[Card 01 Body] FROM: Problem validation, scope management, and overall coordination.

➔ **TO: Problem validation, PRD governance specs, workflow architecture, and submission delivery.**

[Card 02 Title] FROM: AI/ML & Verification

➔ **TO: [Insert Name] — AI & Verification Engine**

[Card 02 Body] FROM: Claim extraction, evidence matching, and rule-based verification.

➔ **TO: Deterministic rule engine, claim extraction, schema validation, and Evidence-or-Abstain pipeline.**

[Card 03 Title] FROM: Backend & Database

➔ **TO: [Insert Name] — Backend & Platform**

[Card 03 Body] FROM: API development, Brand Truth Registry, data storage, and authentication.

➔ **TO: FastAPI services, in-memory vault, D1/Postgres schemas, and role-based permissions.**

[Card 04 Title] FROM: Frontend & UI/UX

➔ **TO: [Insert Name] — Frontend & UX Engineering**

[Card 04 Body] FROM: User interface, content submission, findings display, and review workflow.

➔ **TO: Next.js App Router, responsive audit queue, diff-based correction editor, and review portal.**

[Card 05 Title] FROM: Testing & Documentation

➔ **TO: [Insert Name] — QA & Systems Transparency**

[Card 05 Body] FROM: Workflow testing, reliability checks, documentation, and demo preparation.

➔ **TO: E2E/API test suites, fail-soft error handling, AI/Privacy disclosure docs, and live demo harness.**

SLIDE 10: THANK YOU (CLOSING SLIDE)

- **Title:** THANK YOU
- **Brand Name:** PROOFLINE
- **Subtitle:** Evidence-Backed AI for Brand Governance
- **Tagline:** Every claim checked. Every finding traceable.
- **Recommended Addition:** Live Demo Link: [Insert Live Demo URL]
- **Recommended Addition:** GitHub Code Repository: [Insert GitHub Repo URL]
- **Recommended Addition:** Team / Contact: Binary Brains (potlasrisharanchowdary@gmail.com)